



# Introduction to Deep Learning, PyTorch and TorchPhysics (A3)

## Part I: Introduction to Deep Learning

Daniel Otero Baguer, Janek Gödeke

Center for Industrial Mathematics (ZeTeM)  
University of Bremen

Bremen  
17-21.07.2023





# Outline

- 1 Introduction
- 2 Convolutional Neural Networks
- 3 Architectures
- 4 Data augmentation
- 5 Transfer learning





## Section 1

# Introduction



# Sources

- *“Deep Learning for Vision Systems”*, Mohamed Elgendy  
<https://livebook.manning.com/book/deep-learning-for-vision-systems/welcome/v-8>
- *“Convolutional Neural Networks for Visual Recognition”*  
<http://cs231n.stanford.edu/>

# The learning task

Let's consider a simple regression problem where the observations are:

$$(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)}) \in X \times Y \quad (1)$$

**Task:** Infer a function  $f$  from the training data, which minimizes

$$R(f) = \int_{X \times Y} c(x, y, f(x)) dP(x, y) \quad (2)$$

where  $c$  is a loss function. A common choice is  $c(x, y, f(x)) = \frac{1}{2} \|f(x) - y\|^2$ .

# The learning task

## Empirical Risk

$$R(f) = \frac{1}{m} \sum_{i=1}^m c(x^{(i)}, y^{(i)}, f(x^{(i)})) \quad (3)$$

If we allow  $f$  to be any function that maps from  $X$  to  $Y$ , then we can minimize (5) but at the same time be very distant from the minimizer of (2).

$$f(x) = \begin{cases} y^{(i)} & x = x^{(i)} \\ 0 & \text{otherwise} \end{cases} \quad (4)$$

# The learning task

## No Free lunch theorem


If we do not make any assumptions on the class of functions where  $f$  belongs to, there is no chance to learn anything.



# Learning from data with Neural Networks (NN)

Given observations:  $\{(x^{(m)}, y^{(m)})\}$  and model  $\varphi_{\theta} : X \rightarrow Y$  (NN),  
minimize empirical error:

$$\arg \min_{\theta \in \Theta} = \frac{1}{m} \sum_{i=1}^m c(x^{(i)}, y^{(i)}, \varphi_{\theta}(x^{(i)})) \quad (5)$$







# Learning from data with Neural Networks (NN)

[Published: 09 October 1986](#)

## Learning representations by back-propagating errors


[David E. Rumelhart](#), [Geoffrey E. Hinton](#) & [Ronald J. Williams](#)

[Nature](#) **323**, 533–536 (1986) | [Cite this article](#)

**112k** Accesses | **14906** Citations | **385** Altmetric | [Metrics](#)

# Why Deep learning now?

- Five decades of research in machine learning
- CPUs/GPUs/storage developed for other purposes
- Lots of data from the internet
- Imagenet



# Why Deep learning now?

- Five decades of research in machine learning
- CPUs/GPUs/storage developed for other purposes
- Lots of data from the internet
- Imagenet




# Why Deep learning now?

- Five decades of research in machine learning
- CPUs/GPUs/storage developed for other purposes
- Lots of data from the internet
- Imagenet




# Why Deep learning now?

- Five decades of research in machine learning
- CPUs/GPUs/storage developed for other purposes
- Lots of data from the internet
- Imagenet



# Why Deep Learning now?

## Revolution of Depth



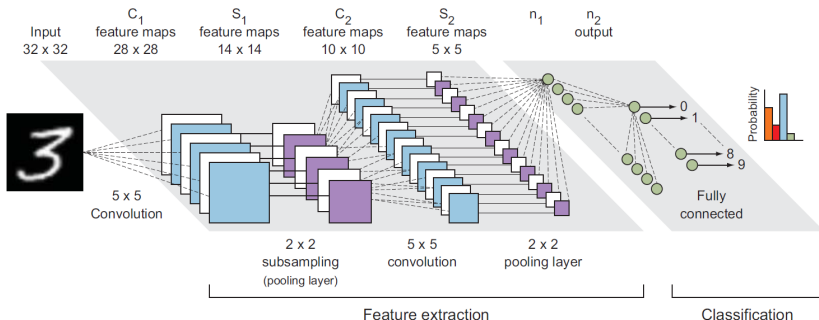


## Section 2

# Convolutional Neural Networks




# Convolutional Neural Network



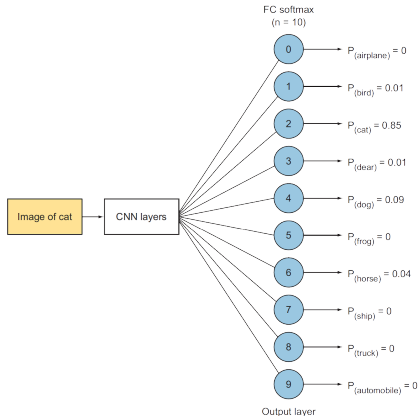


- Input




# Layer types

## ■ Output



# Layer types

## ■ Convolutions



# Layer types

## ■ Convolutions




Figure: Receptive field

# Layer types

## ■ Convolutions

Input image



Edge detection  
kernel

\*

0	-1	0
-1	4	-1
0	-1	0

=

Convolved image  
(feature map)



Figure: Convolutional layers are inspired on standard image filtering

# Layer types

## ■ Pooling





Figure: Max pooling example

# Layer types

## ■ Dropout



### Dropout: A Simple Way to Prevent Neural Networks from Overfitting

*Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, Ruslan Salakhutdinov*, 15(56):1929–1958, 2014.

# Tensors




Figure: Tensors <sup>1</sup>

<sup>1</sup>Image source: <https://towardsdatascience.com/understanding-pytorch-with-an-example-a-step-by-step-tutorial>






## Section 3

# Architectures



# LeNet


- Introduced in 1998, by LeCun et al. in their paper  
“Gradient-Based Learning Applied to Document Recognition”
- 5 layers: 3 conv + 2 fully-connected
- 61706 parameters



# AlexNet

- Winner of the ILSVRC image classification competition in 2012.
- Introduced by Alex Krizhevsky, Geoffrey Hinton and Ilya Sutskever in their paper “ImageNet Classification with Deep Convolutional Neural Networks”
- 8 layers: 5 conv + 3 fully-connected
- 60 million parameters

# AlexNet



# How deep can we go?

- Very deep networks are able to represent very complex functions
- The network can learn features at many different levels of abstraction

## Vanishing gradients

- By the chain rule, the derivatives of each layer are multiplied down the network
- Gradient decreases exponentially as we propagate down to the initial layers
- First hidden layers are learning much slower than later hidden layers

# How deep can we go?




Figure: Sigmoid activation

# How deep can we go?




Figure: Vanishing gradient effect. First layers train much slower.

# How deep can we go?

## Solutions:

- Sigmoid with restricted inputs
- Use ReLu activations
- Normalization layers
- Residual Networks



# How deep can we go?

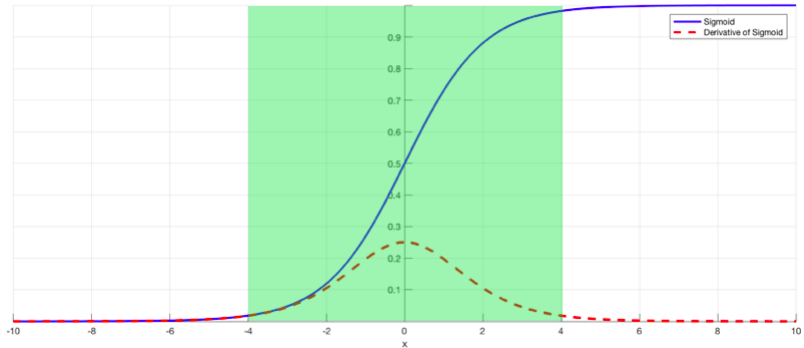


Figure: Sigmoid with restricted inputs

# How deep can we go?




Figure: ReLU (Rectifier linear unit)

---

## Deep Sparse Rectifier Neural Networks

---

Xavier Glorot  
DIRO, Université de Montréal  
Montréal, QC, Canada  
glorotxa@iro.umontreal.ca

Antoine Bordes  
Heudiasyc, UMR CNRS 6599  
UTC, Compiègne, France  
and  
DIRO, Université de Montréal  
Montréal, QC, Canada  
antoine.bordes@hds.utc.fr

Yoshua Bengio  
DIRO, Université de Montréal  
Montréal, QC, Canada  
bengioy@iro.umontreal.ca

# How deep can we go?

## Residual networks (ResNet)

- Introduces shortcut (skip-connection) that allows the gradient to directly back-propagate to earlier layers
- Allows the layer to learn an identity function (will perform at least as good as the previous layer)




Figure: ResNet's skip connection

# How deep can we go?

- From 18 to 152 layers
- With 152 layers:  $\approx 60$  million parameters




Figure: Residual block

# How deep can we go?




Figure: ResNet



## Section 4

# Data augmentation



# Data augmentation

**Main idea:** Generate additional images based on existing ones

- Random flip
- Random erasing
- Random sized crop
- Random perspective
- Color jitter




Figure: Data augmentation

# Data augmentation

## Insights

- The epoch size does not change
- We get different randomly transformed samples every epoch (dynamic dataset)
- Helps to avoid overfitting

## Advise

- Applying heavy augmentations unnecessarily can result in poor accuracy
- Do not use data augmentation in validation or test set



# Data augmentation

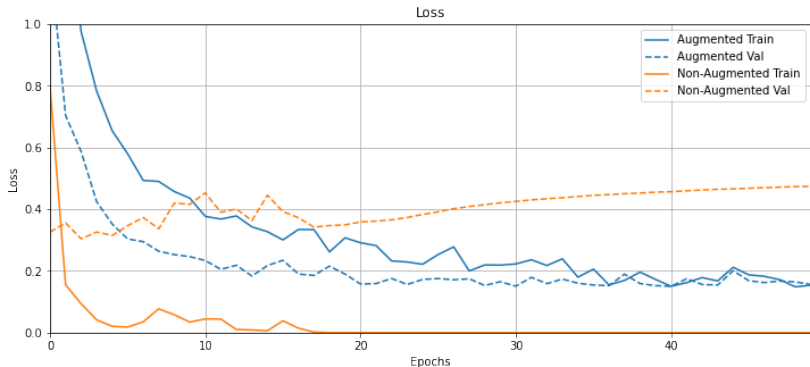


Figure: Non-augmented model vs augmented model <sup>2</sup>.

<sup>2</sup>[https://www.tensorflow.org/tutorials/images/data\\_augmentation](https://www.tensorflow.org/tutorials/images/data_augmentation)



## Section 5

# Transfer learning



# Transfer learning

Common phenomena on CNNs trained on natural images:



**Figure:** First layer learn features similar to Gabor filters and color blobs. Each of the 96 filters shown here is of size  $[11 \times 11 \times 3]$ ,

**Claim:** Such a layer is not specific to a particular dataset or task

# Transfer learning

**Main idea:** Use a pre-trained network

## Use cases

- Fixed feature extractor: Take a pre-trained network and train only the last fully-connected layer
- Fine-tuning: Fine-tune the weights of the pre-trained network by continuing the back-propagation
  - Use small learning rate
  - Keep some of the earlier layers fixed (due to overfitting concerns)

## When and how?

	Small dataset	Large dataset
Similar	Fixed feature extractor + Linear Classifier	Fine-tuning
Very different	Keep early layers fixed + Linear classifier	Pre-trained as initialization

# Practical advice

- Constraints from pre-trained models
- Small learning rate for fine-tuning (avoid distorting pre-trained parameters too quickly or too much)

# Turing award





Thanks!

